

A Comparison of Deep Learning Approaches for Power-based Side-channel Attacks

Roberto Capoferri, Alessandro Barengi, Luca Breveglieri,
Niccolò Izzo, Gerardo Pelosi

November 6th, 2024



POLITECNICO
MILANO 1863

Outline

Introduction and objectives

Background

- Deep learning

- Side-channel attacks

- Deep learning based side-channel attacks

Methodology

- Data collection

- Hyperparameter tuning

- Training and performance evaluation

Performance results

Concluding remarks



POLITECNICO
MILANO 1863

Introduction

Side-channel attacks (SCAs) are a very relevant threat to the devices that offer cryptographic operations, even modern ones [9, 17].

SCAs exploit *information leakage* from the target to extract secret information.

Deep learning applied to SCA: overview

1. collection of *traces* from the target device
2. build the Neural Network (NN) which takes traces as input and produces a probability distribution over the keyspace
 - 2.1 tune the model hyperparameters
 - 2.2 training on the whole dataset: a classification task on target intermediate
3. testing: determine the *Guessing Entropy* (GE) on a different set of traces

Details of all the points in the methodology section.



Objectives

Provide a systematic comparison of the effects of:

- ▶ single/multiple devices to determine *portability*
- ▶ fixed/random key traces
- ▶ different target intermediates
- ▶ usage of plaintext information

Find the combination with the best attack performance.



Outline

Introduction and objectives

Background

- Deep learning

- Side-channel attacks

- Deep learning based side-channel attacks

Methodology

- Data collection

- Hyperparameter tuning

- Training and performance evaluation

Performance results

Concluding remarks




POLITECNICO
MILANO 1863

Background – Deep Learning (DL)

Branch of machine learning that focuses on neural networks.

We consider the Multi-Layer Perceptron (MLP), composed by multiple, fully connected layers of neurons:

- ▶ **input layer:** each neuron takes an input feature
- ▶ one or more **hidden layers:** transform the input features through some weight matrices and non-linear activations
- ▶ **output layer:** the number and the role of each neuron depend on the task



Background – Side-channel attacks (SCAs)

During cryptographic operations, collect *traces* of:

- ▶ execution time
- ▶ power consumption [7]
- ▶ electromagnetic emission [13]

Build a *leakage function* that correlates such measurements with the secret key.

- ▶ **non-profiled SCA**: measurements done on the target device, extract key by using statistical analysis (e.g., DPA [8], CPA [3])
- ▶ **profiled SCA**: collect traces from a clone device to build a model which is then used to attack the target (e.g., template attack [4])



Background – DL applied to SCAs

It's a type of profiled attack, can outperform traditional methods [10] and has some advantages:

- ▶ automatic feature extraction
- ▶ no hypotheses on noise distribution

The profiling phase (a.k.a. *training phase*) can be completely automated.
After the network is trained, only a few traces are needed to perform the attack.



Outline

Introduction and objectives

Background

- Deep learning

- Side-channel attacks

- Deep learning based side-channel attacks

Methodology

- Data collection

- Hyperparameter tuning

- Training and performance evaluation

Performance results

Concluding remarks



POLITECNICO
MILANO 1863

Methodology

Evaluate different approaches to dataset collection and training:

- ▶ **target**: observe the effect of changing the target intermediate
 - ▶ SBOX_OUT: targets the output of the SBox
 - ▶ HW_S0: targets the hamming weight of the SBox output
- ▶ **multi-device**: check portability problem [2] and performance of single-device vs multi-device training in different scenarios
- ▶ **plaintext**: add plaintext information to raw trace data as input to MLP

All code available [16]



Methodology – Data collection

Traces captured on a Riscure Pinata, based on an STM32F4 (Cortex-M4) 32-bit microcontroller running at 168MHz.

Board instrumentation

- ▶ **current probe** to measure power consumption
- ▶ **serial connection** to handle communication
- ▶ **trigger probe** to start the capture on the scope

Oscilloscope: Tektronik MSO58, 625MS/s, 8 bit resolution




Figure: Capture setup



Methodology – Data collection

Traces are captured from **three different devices** running a software implementation of AES128.

Capture SBox computation

1. start AES computation
2. raise GPIO pin before *SBox* operation
3. lower GPIO pin before *MixColumns* operation

Dataset publicly available
for further studies [14,
15]

For each device

- ▶ 200,000 traces collected with a fixed key (key 0)
- ▶ 200,000 traces collected with a random key
- ▶ 30,000 traces collected with a different fixed key (key 1), used for testing



Methodology – Hyperparameter tuning

Explore different MLP structures to get the best model.

Genetic algorithm [5]

Start by selecting random parameters:

1. train model
2. select top performing
3. produce offspring
4. iterate

Table: Hyperparameter space

parameter	possible values
hidden layers	[1, 2, 3, 4, 5]
hidden neurons	[100, 200, 300, 400, 500]
dropout rate	[0.0, 0.1, 0.2, 0.3]
l2	$[0.0, 5 \cdot 10^{-2}, 1 \cdot 10^{-2}, 5 \cdot 10^{-3}, 1 \cdot 10^{-3}, 5 \cdot 10^{-4}, 1 \cdot 10^{-4}]$
optimizer	['adam', 'rmsprop', 'sgd']
learning rate	$[5 \cdot 10^{-3}, 1 \cdot 10^{-3}, 5 \cdot 10^{-4}, 1 \cdot 10^{-4}, 5 \cdot 10^{-5}, 1 \cdot 10^{-5}]$
batch size	[128, 256, 512, 1024]

Table: Parameters for the genetic algorithm

nGen	popSize	selPerc	scProb	mProb
20	15	30%	20%	20%



Methodology – Training and performance evaluation

Use a 90/10 training-validation split and rescale input:

$$X_{\text{scaled}} = (X - \mathbf{min}(X)) / (\mathbf{max}(X) - \mathbf{min}(X))$$

Use of regularization options to improve model performance

- ▶ **Dropout:** which percentage of neurons to turn off in a layer during training.
- ▶ **L2 regularization:** limits the value of the weights by adding a penalty to larger values.
- ▶ **Batch normalization** [6]: normalizes the input to each layer, stabilizing the training process.
- ▶ **Early stopping** [12]: monitors training and validation accuracy, and stops the training process if overfitting is detected

Evaluation metric: Guessing Entropy [18]

- ▶ rank of the correct key byte over an increasing number of test traces
- ▶ average over multiple runs



Outline

Introduction and objectives

Background

- Deep learning

- Side-channel attacks

- Deep learning based side-channel attacks

Methodology

- Data collection

- Hyperparameter tuning

- Training and performance evaluation

Performance results

Concluding remarks



POLITECNICO
MILANO 1863

Results

Performance

- ▶ measured using the GE
- ▶ graphs show evolution of the position of the correct key when incrementing the number of attack traces
- ▶ tested from 1 to 300 attack traces
- ▶ average over 100 runs on different traces (use up to all 30,000 test traces)

Scenarios

- ▶ attack the same device(s) as the training one(s)
- ▶ attack a different device




Figure: Fixed key, 1 device w/ ptx, HW_S0. Vertical lines where $GE < 0.5$



Results – Plaintext

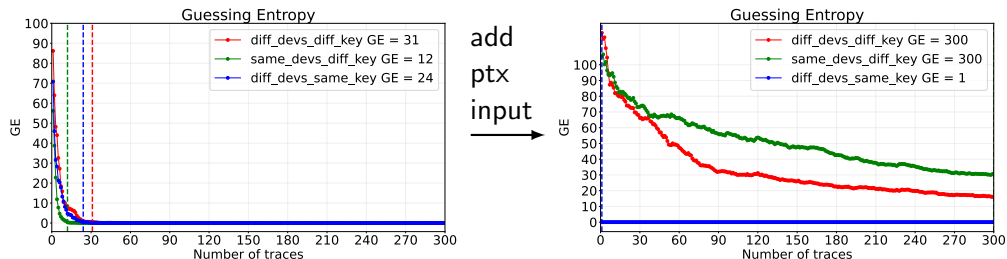


Figure: Effect of ptx on model performance using SBOX_OUT

Fixed key scenario

- ▶ adding plaintext information breaks the model when considering SBOX_OUT
- ▶ for HW_S0 the model still works but performance is works
- ▶ probably due to overfitting on the key



Results – Plaintext





Figure: Effect of ptx on model performance using SBOX_OUT

Random key scenario

- ▶ in this case performance is improved
- ▶ similar behaviour with both SBOX_OUT and HW_S0 targets
- ▶ also applies when considering multiple training devices




Results – Multi-device



Fixed key GE with 1 training device.



Results – Multi-device

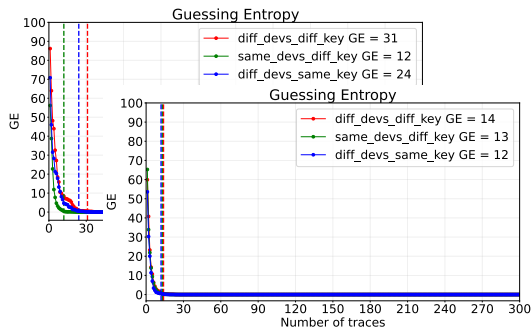


Multi-device

- ▶ improves performance in general
- ▶ reduces performance loss when considering portability




Results – Multi-device



Multi-device


- ▶ improves performance in general
- ▶ reduces performance loss when considering portability



Random key GE with 1 training device.





Results – Multi-device



Multi-device

- ▶ improves performance in general
- ▶ reduces performance loss when considering portability
- ▶ has a greater impact on random key



Random key GE with 1 and 2 training devices.



Results – Target intermediate

- ▶ different behaviour depending on target
- ▶ SBOX_OUT performs better overall in different situations
- ▶ there are cases when HW_S0 is better (e.g., fixed key dataset w/ ptx)

Reason

Using the HW_S0 model can induce a class imbalance [11] in the dataset, impacting performance.



Results – Recap

Effect of different parameters:

Plaintext

- ▶ for fixed key datasets it breaks the model
- ▶ performance increase for random key, but only in SBOX_OUT scenario

Data collection

- ▶ fixed key works better than random key
- ▶ random key can get closer when adding plaintext information

Multi-device and portability

- ▶ confirm that model performs worse when tested on new device
- ▶ using traces from multiple devices yields a more robust model
- ▶ some exception can happen, random key in the HW_SO scenario perform worse



Outline

Introduction and objectives

Background

- Deep learning

- Side-channel attacks

- Deep learning based side-channel attacks

Methodology

- Data collection

- Hyperparameter tuning

- Training and performance evaluation

Performance results

Concluding remarks



POLITECNICO
MILANO 1863

Concluding remarks

Achieved results

- ▶ studied the effect of different parameters for training
- ▶ new study on plaintext information
- ▶ validate multi-device model for portability
- ▶ provide a new dataset to further study portability

Future work

- ▶ different implementations
- ▶ different points of interest
- ▶ different CPU architectures [1]



Questions?



POLITECNICO
MILANO 1863

References I

- [1] Alessandro Barengi and Gerardo Pelosi. "Side-channel security of superscalar CPUs: evaluating the impact of micro-architectural features". In: *Proceedings of the 55th Annual Design Automation Conference, DAC 2018, San Francisco, CA, USA, June 24-29, 2018*. ACM, 2018, 120:1–120:6. DOI: 10.1145/3195970.3196112. URL: <https://doi.org/10.1145/3195970.3196112>.
- [2] Shivam Bhasin et al. "Mind the Portability: A Warriors Guide through Realistic Profiled Side-channel Analysis". In: *27th Annual Network and Distributed System Security Symposium, NDSS 2020, San Diego, California, USA, February 23-26, 2020*. The Internet Society, 2020. URL: <https://www.ndss-symposium.org/ndss-paper/mind-the-portability-a-warriors-guide-through-realistic-profiled-side-channel-analysis/>.
- [3] Eric Brier, Christophe Clavier, and Francis Olivier. "Optimal Statistical Power Analysis". In: *IACR Cryptol. ePrint Arch.* (2003), p. 152. URL: <http://eprint.iacr.org/2003/152>.
- [4] Suresh Chari, Josyula R. Rao, and Pankaj Rohatgi. "Template Attacks". In: *Cryptographic Hardware and Embedded Systems - CHES 2002, 4th International Workshop, Redwood Shores, CA, USA, August 13-15, 2002, Revised Papers*. Ed. by Burton S. Kaliski Jr., Çetin Kaya Koç, and Christof Paar. Vol. 2523. Lecture Notes in Computer Science. Springer, 2002, pp. 13–28. DOI: 10.1007/3-540-36400-5_3. URL: https://doi.org/10.1007/3-540-36400-5_3.
- [5] Matt Harvey and Norman Heckscher. *Evolve a neural network with a genetic algorithm*. <https://github.com/harvitronix/neural-network-genetic-algorithm>. 2017.
- [6] Sergey Ioffe and Christian Szegedy. "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift". In: *Proceedings of the 32nd International Conference on Machine Learning, ICML 2015, Lille, France, 6-11 July 2015*. Ed. by Francis R. Bach and David M. Blei. Vol. 37. JMLR Workshop and Conference Proceedings. JMLR.org, 2015, pp. 448–456. URL: <http://proceedings.mlr.press/v37/ioffe15.html>.
- [7] Paul C. Kocher, Joshua Jaffe, and Benjamin Jun. "Differential Power Analysis". In: *Advances in Cryptology - CRYPTO '99, 19th Annual International Cryptology Conference, Santa Barbara, California, USA, August 15-19, 1999, Proceedings*. Ed. by Michael J. Wiener. Vol. 1666. Lecture Notes in Computer Science. Springer, 1999, pp. 388–397. DOI: 10.1007/3-540-48405-1_25. URL: https://doi.org/10.1007/3-540-48405-1_25.
- [8] Paul C. Kocher et al. "Introduction to differential power analysis". In: *J. Cryptogr. Eng.* 1.1 (2011), pp. 5–27. DOI: 10.1007/S13389-011-0006-Y. URL: <https://doi.org/10.1007/s13389-011-0006-y>.
- [9] Oleksiy Lisovets et al. "Let's Take it Offline: Boosting Brute-Force Attacks on iPhone's User Authentication through SCA". In: *IACR Trans. Cryptogr. Hardw. Embed. Syst.* 2021.3 (2021), pp. 496–519. DOI: 10.46586/TCHES.V2021.I3.496-519. URL: <https://doi.org/10.46586/tches.v2021.i3.496-519>.

References II

- [10] Houssem Maghrebi, Thibault Portigliatti, and Emmanuel Prouff. "Breaking Cryptographic Implementations Using Deep Learning Techniques". In: *Security, Privacy, and Applied Cryptography Engineering - 6th International Conference, SPACE 2016, Hyderabad, India, December 14-18, 2016, Proceedings*. Ed. by Claude Carlet, M. Anwar Hasan, and Vishal Saraswat. Vol. 10076. Lecture Notes in Computer Science. Springer, 2016, pp. 3–26. DOI: 10.1007/978-3-319-49445-6_1. URL: https://doi.org/10.1007/978-3-319-49445-6%5C_1.
- [11] Stjepan Picek et al. "The Curse of Class Imbalance and Conflicting Metrics with Machine Learning for Side-channel Evaluations". In: *IACR Trans. Cryptogr. Hardw. Embed. Syst.* 2019.1 (2019), pp. 209–237. DOI: 10.13154/TCHES.V2019.I1.209-237. URL: <https://doi.org/10.13154/tches.v2019.i1.209-237>.
- [12] Lutz Prechelt. "Early Stopping - But When?" In: *Neural Networks: Tricks of the Trade - Second Edition*. Ed. by Grégoire Montavon, Genevieve B. Orr, and Klaus-Robert Müller. Vol. 7700. Lecture Notes in Computer Science. Springer, 2012, pp. 53–67. DOI: 10.1007/978-3-642-35289-8_5. URL: https://doi.org/10.1007/978-3-642-35289-8%5C_5.
- [13] Jean-Jacques Quisquater and David Samyde. "ElectroMagnetic Analysis (EMA): Measures and Counter-Measures for Smart Cards". In: *Smart Card Programming and Security, International Conference on Research in Smart Cards, E-smart 2001, Cannes, France, September 19-21, 2001, Proceedings*. Ed. by Isabelle Attali and Thomas P. Jensen. Vol. 2140. Lecture Notes in Computer Science. Springer, 2001, pp. 200–210. DOI: 10.1007/3-540-45418-7_17. URL: https://doi.org/10.1007/3-540-45418-7%5C_17.
- [14] Roberto Capoferri. *Fixed-key dataset for three riscure Pinata devices*. <https://zenodo.org/records/11443025>. 2024.
- [15] Roberto Capoferri. *Random-key dataset for three riscure Pinata devices*. <https://zenodo.org/records/11199202>. 2024.
- [16] Roberto Capoferri. *Source code and models*. <https://github.com/RobertoCapoferri/DLSCA-article>. 2024.
- [17] Thomas Roche. *EUCLEAK*. Cryptology ePrint Archive, Paper 2024/1380. 2024. URL: <https://eprint.iacr.org/2024/1380>.
- [18] Lichao Wu et al. *On the Attack Evaluation and the Generalization Ability in Profiling Side-channel Analysis*. Cryptology ePrint Archive, Paper 2020/899. 2020. URL: <https://eprint.iacr.org/2020/899>.

Appendix A - full comparison tables I

Table: Results for the **fixed key** dataset. No. of traces to get $GE < 0.5$

target	variant	n. devices	same_devs_diff_key	diff_devs_same_key	diff_devs_diff_key
SBOX OUT	no ptx	1	12	24	31
		2	13	12	14
	ptx	1	> 300	1	> 300
		2	> 300	1	> 300
HW SO	no ptx	1	20	29	35
		2	18	13	28
	ptx	1	34	18	68
		2	> 300	8	> 300

Appendix A - full comparison tables II

Table: Results for the **random key** dataset. No. of traces to get $GE < 0.5$

target	variant	n. devices	same_devs_key_1	diff_devs_key_0	diff_devs_key_1
SBOX OUT	no ptx	1	249	143	> 300
		2	45	51	70
	ptx	1	39	42	84
		2	20	25	25
HW SO	no ptx	1	165	> 300	> 300
		2	117	70	117
	ptx	1	75	192	188
		2	> 300	> 300	> 300